

facebook

cgroupv2: Linux's new unified control group system

Chris Down (cdown@fb.com)

Production Engineer, Web Foundation

In this talk

- Short intro to cgroups
- v1, what went well, what didn't
- Design decisions in v2
- New features
- Future work

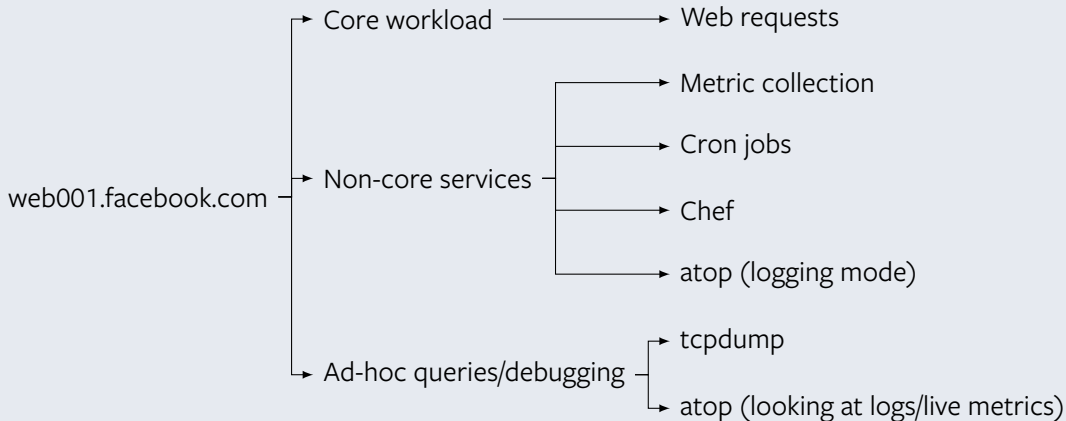
About me

- Chris Down
- Production Engineer, Web Foundation
- Working at Facebook London
- `cdown@fb.com`


Why WF cares about cgroupv2

- >100,000 servers
- Many thousands of services
- Want to limit failure domains

The problem



Limits




What are cgroups?

- `cgroup` \equiv control group
- System for resource management on Linux
- Directory hierarchy at `/sys/fs/cgroup`
- Limit, throttle, and account for resource usage per control group
- Each resource interface is provided by a *controller*

Practical uses

- Isolating core workload from background resource needs
 - Web server vs. system processes (eg. Chef, metric collection, etc)
 - Time critical work vs. long-term asynchronous jobs
- Don't allow one workload to overpower the others

Who uses cgroups?



How did this work in cgroupv1?


cgroupv1 has a hierarchy per-resource, for example:

```
% ls /sys/fs/cgroup
cpu/  cpuacct/  cpuset/  devices/  freezer/
memory/  net_cls/  pids/
```


Each resource hierarchy contains cgroups for this resource:

```
% find /sys/fs/cgroup/pids -type d
/sys/fs/cgroup/pids/background.slice
/sys/fs/cgroup/pids/background.slice/async.slice
/sys/fs/cgroup/pids/workload.slice
```


How did this work in cgroupv1?




How did this work in cgroupv1?



How did this work in cgroupv1?



Hierarchy in cgroupv1



How does this work in cgroupv2?


cgroupv2 has a *unified hierarchy*, for example:

```
% ls /sys/fs/cgroup
background.slice/  workload.slice/
```

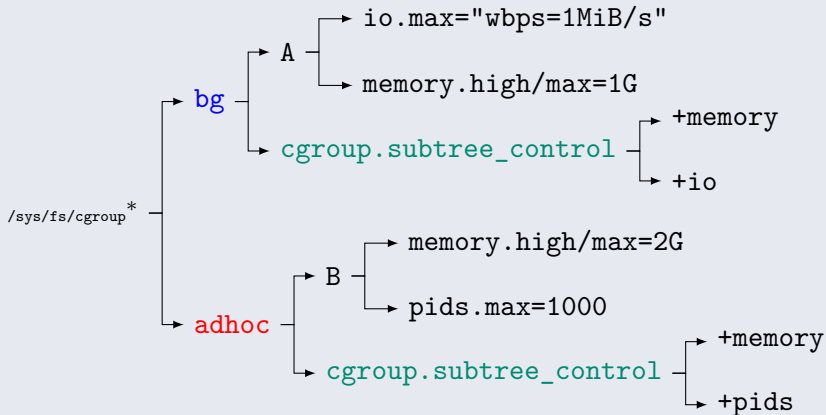
Each cgroup can support multiple resource domains:

```
% ls /sys/fs/cgroup/background.slice
async.slice/  foo.mount/  cgroup.subtree_control
memory.high  memory.max  pids.current  pids.max
```

How does this work in cgroupv2?




Hierarchy in cgroupv2




* in real life, we must enable memory/pids/io controllers for children here too

Hierarchy in cgroupv1



Hierarchy in cgroupv2





* in real life, we must enable memory/pids/io controllers for children here too

Fundamental differences between v1 and v2

- Unified hierarchy — resources apply to cgroups now
- Granularity at TGID (PID), not TID level
- Focus on simplicity/clarity over ultimate flexibility

“No internal process” constraint





Practical improvements in v2

v2 improvements: Tracking of non-immediate/multi-resource charges

In v1:

- No tracking of non-immediate charges
- Charged to root cgroup, essentially unlimited

In v2:

- Page cache writebacks and network are charged to the responsible cgroup
- Can be considered as part of cgroup limits

v2 improvements: Communication with backing subsystems

In v1:

- Most actions for non-share based resources reacted crudely to hitting thresholds
- For example, in the memory cgroup, the only option was to OOM kill or freeze*

In v2:

- Many cgroup controllers negotiate with subsystems before real problems occur
- Subsystems can take remediative action (eg. direct reclaim with `memory.high`)
- Easier to deal with temporary spikes in a resource's usage

* `soft_limit_in_bytes` exists in v1, but it's too overloaded and abstract to be useful

v2 improvements: Saner notifications

In v1:

- One `clone()` per event for cgroup release, expensive
- `eventfd()` support for others

In v2:

- inotify support everywhere
- `eventfd()` support still exists
- One process to monitor everything, if you like

v2 improvements: Utility controllers make sense now

In v1:

- Utility controllers have their own hierarchies
- We usually want to use processes from another hierarchy
- As such, we end up manually synchronising 😞

In v2:

- We have a single unified hierarchy, so no sync needed

v2 improvements: Consistency between controllers

In v1:

- Inconsistencies in controller APIs
- Some controllers don't inherit values

In v2:

- Our crack team of API Design Experts have ensured your sheer delight*

* well, at least we can pray we didn't screw it up too badly

v2 improvements: Unified limits

In v1:

- Some limitations could not be fixed due to backwards compatibility
- `memory.{,kmem.,kmem.tcp.,memsw.,[...] }limit_in_bytes`

In v2:

- Less iterative, more designed up front
- We now have universal thresholds (eg. `memory.{high,max}`)


Use at FB

- ~10% of web servers
- Already getting better results on spiky workloads
- Separation of workload services from system services (eg. Chef, metric collection)
- Running managed with systemd (see “Deploying systemd at scale” talk)

Current support

- Merged:
 - I/O
 - Memory
 - PID
- Complete, but pending merge:
 - CPU (thread: <https://bit.ly/cgroupv2cpu>)

v2 big bets: Real memory pressure detection



Future work

- CPU accounting for “tracked” events (eg. page cache writeback)
- Work-conserving I/O controller
- Accounting for filesystem metadata updates
- Better metrics around memory pressure (eg. refault metrics)
- Freezer for v2

How can I get it?

- `systemd.unified_cgroup_hierarchy=1`
 - systemd will mount `/sys/fs/cgroup` as cgroupv2
 - Available from systemd v226 onwards
- `cgroup_no_v1=all`
 - The kernel will disable all v1 cgroup controllers
 - Available from Linux 4.6 onwards

Manual mounting: `% mount -t cgroup2 none /sys/fs/cgroup`

Links

- Docs at <https://bit.ly/cgroupv2>
- Come ask me questions 😊

facebook